

Event-Object Reasoning with Curated Knowledge Bases: Deriving Missing Information

Chitta Baral and Nguyen H. Vo

School of Computing, Informatics, and Decision Systems Engineering,
Arizona State University, Tempe, Arizona, USA

Abstract. The broader goal of our research is to formulate answers to why and how questions with respect to knowledge bases, such as AURA. One issue we face when reasoning with many available knowledge bases is that at times needed information is missing. Examples of this include partially missing information about next sub-event, first sub-event, last sub-event, result of an event, input to an event, destination of an event, and raw material involved in an event. In many cases one can recover part of the missing knowledge through reasoning. In this paper we give a formal definition about how such missing information can be recovered and then give an ASP implementation of it. We then discuss the implication of this with respect to answering why and how questions.

1 Introduction

Our work in this paper is part of two related long terms goals: answering “How”, “Why” and “What-if” questions and reasoning with the growing body of available knowledge bases ¹, some of which are crowd-sourced. Although answering “How” and “Why” questions are important, so far little research has been done on them. Our starting point to address them has been to formulate answers to such questions with respect to abstract knowledge structures obtained from knowledge bases. In particular, in the recent past we considered Event Description Graphs (EDGs) [1] and Event-Object Description Graphs (EODGs) [2] to formulate answers to some “How” and “Why” questions with respect to the Biology knowledge base AURA [3].

Going from the abstract structures to reasoning with real knowledge bases (KBs) we noticed that the KBs often have missing pieces of information, such as properties of an instance (of a class) or relations between two instances. For example, AURA does not encode that *Eukaryotic translation* is the next event of *Synthesis of RNA in eukaryote*; this may be because the two subevents of “Protein synthesis” were encoded independently. The missing pieces make the KB and the Description Graphs constructed from it fragmented and as a result answers obtained with respect to them are not intuitive. Moreover, the KBs like AURA often have two or more names that refer to the same entity. To get intuitive answers they need to be resolved and merged into a single entity.

¹ See for example, <http://linkeddata.org/>.

Such finding of non-identical duplicates in the KB and merging them into one is referred in the literature as entity resolution [4, 5].

In this paper, we start with introducing knowledge description graphs (KDGs) as structures that can be (without much reasoning) obtained from frame based KBs such as AURA. We discuss underspecified knowledge description graphs (UDGs) and formulate notions of reasoning with respect to these graphs to obtain certain missing information. We then present our approach of entity resolution and use it in recovering additional missing information. We give an Answer Set Programming (ASP) encoding of our formulation. We conclude with a discussion on the use of the above in answering “why” and “how” questions.

2 Background: Frame-based Knowledge Bases; ASP

The KB we used in this work is based on AURA [3] and was described in details in [6]. AURA is a frame-based KB manually curated by biology experts; it contains a large amount of frames describing biological entities events (or processes). One important aspect of our KB is the class hierarchy. For example ²: its basic class is *Thing*, which has two children classes: *Entity* and *Event*. *Entity* is the ancestor of all classes of biological entities; *Event*, of biological events. For instance, *Spatial entity*, *Eukaryote*, *Nucleus* and *mRNA* are descendants of *Entity*, while “Eukaryotic translation”, “Eukaryotic transcription” are descendant of *Event*.

Our KB is a set of facts of the form “has(*A*, slot_name, *B*)” where *A* and *B* are either classes or instances (of classes), *slot_name* is the name of the relation between *A* and *B* such as *instance_of*, *raw_material* or *results*. The statement “*eukaryotic translation* is based on *mRNA*” is represented in our KB as follows.

```
has(euka_transl4191, instance_of, event).
has(euka_transl4191, instance_of, eukaryotic_translation).
has(euka_transl4191, base, mrna4642).
has(mrna4642, instance_of, mrna).
```

This snippet reads as “*eukaryotic_translation4191* is an instance of class *event* and an instance of class *eukaryotic_translation*. *eukaryotic_translation4191* is based on *mrna4642*, which is an instance of *mrna*”.

For the declarative implementation of our formulations, we use ASP [7]. That allows us to use our earlier work [6] on using ASP to reason with frame-based knowledge bases. ASP’s strong theoretical foundation [8] and its default negation and recursion are useful in our encoding and in proving results about them.

3 Knowledge Description Graphs

An **Underspecified Knowledge Description Graph (UDG)** is a structure to represent the facts about instances and classes of events, entities and relationships between them. An UDG is constructed from knowledge bases such as AURA. Formal definition of the UDGs is given in the following.

Definition 1. *An UDG is a directed graph with one type of node and five types of directed edges: compositional edges, ordering edges, class edges, locational edges and participant edges. Each node represents an instance (of a class) or a class in our KBs.*

² Our examples are either directly from AURA, or are slightly modified from it.

Edge type	Relation(s)
locational	happenings
class	instance-of, super-class
compositional	subevent, first-subevent, has-part, has-region, has-basic-structural-unit
ordering	next-event, enables, causes, prevents, inhibits
participant	raw-materials, result, agent, destination, instrument, origin, site

Table 1. Types of edges in an UDG. An edge of relation Y from a node X to Z represents $X[Y] = Z$, meaning the slot Y of the entity X has value Z .

We used the slot names in KM [9] and AURA as a guide to categorize four types of edges (Table 1).

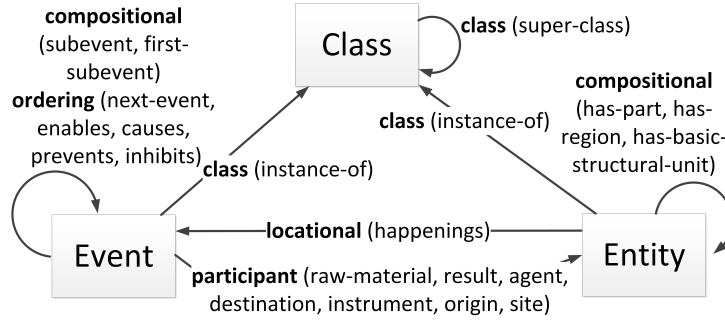


Fig. 1. Types of edges in a KDG

A **Knowledge Description Graph (KDG)** (a slight generalization of EODGs in [2]) is constructed from an UDG. A node in a KDG represents either an instance of a biological entity, an instance of a biological process, or a class of biological entity/event. The KDG structure allows us to answer “How” and “Why” questions. More formally, a KDG is defined as follows.

Definition 2. A KDG is a directed graph with: (i) three types of nodes: event nodes, entity nodes, and class nodes; and (ii) five types of directed edges: compositional edges, class edges, ordering edges, locational edges and participant edges. A KDG has the property that there are no directed cycles within any combination of compositional, locational and participant edges.

Figure 1 shows the types of edges in a KDG and the corresponding sources and destinations of the edges. Edges in a KDG are from the edges of the UDG, with additional type constraints of the source and destination nodes. For example, ordering edges must be from events to events; compositional edges are from events to events or from entity to entity, depending on their specific relations.

Since UDGs and KDGs can be huge, we usually work on their smaller subgraphs that are rooted at an entity or an event. They are defined as follows.

Definition 3. Let Z be a node in a KDG G . The Knowledge Description Graph (KDG) rooted at Z is the subgraph of G composed of: (1) The set N of all the nodes of G that are accessible from Z through compositional edges, class edges,

In other words, E' is the next event of E if there is an ordering edge from E to E' .

Definition 6. Let S be the set of subevents of an event X in the $KDG(Z)$. Event E in S is the first subevent of X if there exists no other event E' in S such that E is the next event of E' . Similarly, event E in S is the last subevent of X if there exists no other event E' in S such that E' is the next event of E .

Here we assume that S was properly encoded in that there is only one chain of subevents in S . In our KB, *light reaction* and *calvin cycle* are two subevents of *photosynthesis* and *light reaction* enables *calvin cycle*. But their orders are not defined. However, using Definition 5 and 6, we can identify that: *calvin cycle* is the next event of *light reaction*; *light reaction* is the first subevent of *photosynthesis*; and *calvin cycle* is the last subevent.

4.2 Input/Output of Events

Two types of events: In our KB there are two types of events: transport events and operational events. In a transport event, there is only a change in the locations; the input location and output location are different from each other while the input entity and output entity are the same. All other events are operational events. In an operational event, there is usually no change in its location. We differentiate two types of events by their ancestor classes; transport events are descendants of the classes *move_through*, *move_into* and *move_out_of*.

Input, Output, Input Location, Output Location: To reason about the KDG, we need the input and output of each event as well as the input location and the output location, which are not always available. In the following, we show how to use various event's relations - such as *raw-material*, *destination*, *location* and others - to create four new relations (IO relations): *input*, *output*, *input-location* and *output-location*. After that, we propose rules to complete the KDG's IO relations.

We created the IO relations of an event based on specific relations as shown in Table 2. The meaning of relation "base" from AURA depends on the context. For transport events, it is for *input-location*; for operational events, it is for *input*.

Event type	IO relation type	Relation(s)
Transport event	input	object
Transport event	output	object
Transport event	input-location	base, origin
Transport event	output-location	destination
Operational event	input	object, base, raw-material
Operational event	output	result
Operational event	input-location	site
Operational event	output-location	destination ³

Table 2. The IO properties of events and their corresponding relations.

Completing Missing Information of Input, Output, Input Location, Output Location: We can obtain missing IO properties of an event from its subevent(s). For instance, an input of the first subevent of E is also an input of E .

Definition 7. Let FSE and LSE respectively be the first subevent and last subevent of event E in $KDG(Z)$.

Let $InputRelation$ be the input relation, input-location relation or one of their corresponding relations (Table 2). If $InputRelation$ is a relation from FSE to X then $InputRelation$ is also a relation from E to X .

Let $OutputRelation$ be the output relation and output-location relation or one of their corresponding relations. If $OutputRelation$ is a relation from LSE to X then $OutputRelation$ is also a relation from E to X .

In our KB, *photosynthesis* has two subevents: *light reaction* and *calvin cycle*, the next event of *light reaction*. *Sunlight* is the raw-material of the *light reaction*, *sugar* is the result of *calvin cycle*. Using Definition 7, we have that *sunlight* is the raw-material of *photosynthesis* and *sugar* is its result. Moreover, we also have: *sunlight* is the input of *light reaction* as well as *photosynthesis*; *sugar* is the output of both *calvin cycle* and *photosynthesis*.

Similarly, the output location of an operational event is often not defined in the KB but we can use input location as the default value for output location.

Definition 8. Let E be an event in $KDG(Z)$, E 's input location is also the output location if E 's output location has not been specified.

Figure 3 shows the IO properties of events in Fig. 2. The properties in bold are the ones that were recovered using Definitions 7 and 8.

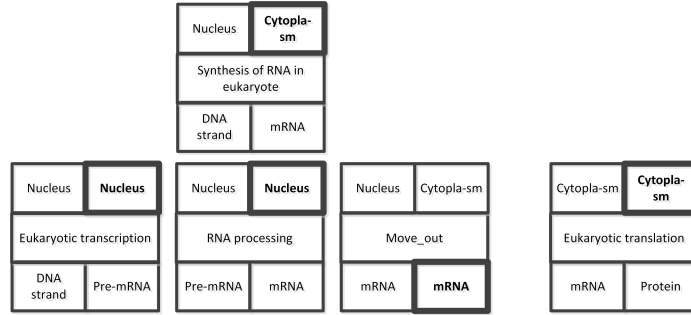


Fig. 3. The IO properties of events in Fig. 2. The five blocks contain IO properties of events: *Synthesis of RNA in eukaryote*, *Eukaryotic translation*, *Move_out*, and *Eukaryotic transcription*. The middle rectangle of each block contains the event name. The top rectangles are for input and output locations; the bottom rectangles are for input and output. The properties in bold were recovered using Definitions 7 and 8

4.3 Main Class of an Instance

In our KB, one instance can belong to many classes. For example, *dna_strand19497* - the input of *Eukaryotic transcription* - is an instance of *dna_strand*, *dna_sequence*, *nucleic_acid* and *polymer*⁴. However, to reason about the equality between in-

⁴ For the sake of simplicity, in the previous figures and descriptions, we usually referenced the entities and events by their “main” class(es) and not by the instances’ names although our KB and our implementation works on instances’ names.

stances, we need the “main” class(es), the most specific class(es) of that instance. Our formal definition of “main” class is given below.

Definition 9. Let E be an instance in $KDG(Z)$. $ClassB$ is a main class of instance E if (1) it is a class of E and (2) it is not the case that there is a $ClassA$ which is a class of E and (a) $ClassB$ is ancestor of $ClassA$ or (b) $ClassB$ is a general class but $ClassA$ is not; where general classes in our KB are *thing*, *event*, *entity*, *spatial_entity*, *tangible_entity*, and *chemical_entity*.

The main classes of *dna_strand19497*, according to the Definition 9, are *dna_strand* and *dna_sequence*; the other classes of *dna_strand19497* are ancestors of those two.

5 Entity Resolution

In the KBs such as AURA, the curation was done in many sessions and probably by many people. (Same is true with respect to many other KBs; especially the ones that are developed using crowd-sourcing.) The results are, in many cases, (i) two different instance names were used when they are probably the same instance; and (ii) parts of some biological process were encoded as independent events. For example: the input of *Eukaryotic translation* (Figure 2) is *mrna4642* whereas the output of *Move_out* is *mrna22911*; *Synthesis of RNA in eukaryote* and *Eukaryotic translation* should be subevents of “Synthesis of protein in eukaryote” but they are encoded as two separate events.

In this section, we propose methods to solve the first problem. These methods are then used to solve the second problem in the next section. In order to compare two instances in a KB, we define a match relation. Generally speaking, instance A can match with instance B if A can be safely used in a context where a term of B is expected. We defined matching relation with many confidence levels for greater flexibility in future works.

Definition 10. Let A and B be two instances in $KDG(Z)$. Let $ClassA$ and $ClassB$ be main classes of A and B respectively.

1. A matches with B with high confidence if one of the following is true
 - (a) $A = B$ (A and B are the same instance)
 - (b) A is cloned from B (Shortcut in AURA to specify that A has all the properties of B)
 - (c) $ClassA$ is an ancestor of $ClassB$.
2. A matches with B with medium confidence if A and B are both cloned from an instance C .
3. A matches with B with low confidence if $ClassA = ClassB$ (A and B are instances of the same main class).
4. A matches with B with confidence $\min(Conf_1, Conf_2)$ if
 - (a) A matches with C with confidence $Conf_1$ and
 - (b) C matches with B with confidence $Conf_2$
5. Otherwise, A does not match with B .

Using Def.10, we can match *mrna4642* - the input of *Eukaryotic translation* - with *mrna22911* - the output of *Move.out*, because both have *mrna* as the main class.

While Def.10 can match all the input and output in our aforementioned example, it is not sufficient for matching location. For example, we can not match an instance of *cytoplasm* to an instance of *cytosol*. However when we say *Event A occurs in cytosol*, we can understand that *Event A occurs in cytoplasm*. To overcome this shortcoming, we define the relation *Spatially match* as follows.

Definition 11. *Instance A in $KDG(Z)$ is a location instance if the class $ClassA$ of A is a descendant of the class *spatial_entity*.*

Definition 12. *Let A and B be two location instances in $KDG(Z)$. Let $ClassA$ and $ClassB$ be main classes of A and B respectively.*

1. *Location A spatially matches with location B with confidence Conf if instance A matches with instance B with confidence Conf.*
2. *Location A spatially matches with location B with high confidence if one of the following is true:*
 - (a) *B is inside A (the relation inside is encoded in our KB by slot name is_inside).*
 - (b) *B is part of A (the relation “part of” is encoded in our KB by slot name part_of).*
3. *Location A spatially matches with location B with confidence $\min(Conf_1, Conf_2)$ if*
 - (a) *A spatially matches with location C with confidence $Conf_1$, and*
 - (b) *C spatially matches with B with confidence $Conf_2$.*

Suppose that in our KB we have: *cytosol234* and *cytosol987* all are instances of *cytosol*; *cytoplasm322* is an instance of *cytoplasm* and *cytosol987* is inside *cytoplasm322*. We can then conclude: *cytosol234* and *cytosol987* match with each other with low confidence, according to Def.10.3; *cytoplasm322* spatially matches with *cytosol987* with high confidence (Def.12.2.a); *cytosol987* spatially matches with *cytosol234* with low confidence (Def.12.1); and *cytoplasm322* spatially matches with *cytosol234* with low confidence (Def.12.3).

6 Finding the Possible Next Events

In this section, we demonstrate the usefulness of matching instances (Definition 10 and 12) in finding the possible next event(s) of a given event. While in simple cases (Section 4.1), we can find the next event E' of an event E by using Definition 5, there are still cases where there exists no ordering edges from E to E' . For examples, *Alteration of mrna ends* and *RNA splicing* are two subevents of *RNA processing* but no other relation between them was defined. However, they all occur in *nucleus16421* and *Alteration of mrna ends*'s output, *pre-mrna7690*, matches with *RNA splicing*'s input, *rna8697*. This information hints us that *RNA splicing* is *Alteration of mrna ends*'s next event.

Following this intuition, our approach for finding the next event is that E' is the next event of E if the output of E matches the input of E' and output

location of E matches the input location of E' . In the example in Figure 2, this assumption holds in all three events: *Eukaryotic transcription*, *RNA processing* and *Move.out*; all of which are already defined in our KB as consequent events. This assumption also suggests that *Eukaryotic translation* can be the next event of either *Synthesis of RNA in eukaryote* or *Move.out*. Armed with Definition 10 and 12, we define the following join relation.

Definition 13. Let A and B be two events in $KDG(Z)$. Event A joins to event B if all of the following conditions are true:

1. The output of A matches with the input of B or vice versa.
2. The output location of A spatially matches with the input location of B or vice versa.

Applying this definition, we have: *Alteration of mrna ends* joins to *RNA splicing*, *Eukaryotic transcription* joins to *RNA processing*; *RNA processing* joins to *Move.out*; and both *Synthesis of RNA in eukaryote* and *Move.out* are joined from *Eukaryotic translation*. Since we want *Eukaryotic translation* to be a possible next event of *Synthesis of RNA in eukaryote* instead of its subevent *Move.out*, we define the possible next event as follows.

Definition 14. Let A and B be two events in $KDG(Z)$ where A joins to B . B is a possible next event of A if none of the following conditions is true:

1. A joins to *Ancestor* B where *Ancestor* B is the ancestor event of B (in other words, there is a non-empty path of subevent relation from *Ancestor* B to B).
2. Ancestor event *Ancestor* A of A joins to B .
3. A is an ancestor event of B .
4. B is an ancestor event of A .
5. A and B have the same ancestor event.

In our example, condition 14.2 gives us that *Eukaryotic translation* is not the possible next event of *Move.out* while 14 concludes that *Eukaryotic translation* is the possible next event of *Synthesis of RNA in eukaryote*. We assume that an event and its subevents are put in our KB as a whole, so the *next_event* relations between them are well defined. Thus conditions 14.3-5 take those relations out of consideration.

When we have a path of possible next events, we can create an event SE , which is the super event of all events in the path, and add suitable *next_event* or *subevent* relations. This new event would link the events that were mistakenly encoded as independent events (that we mentioned earlier).

7 ASP Encodings

In this section we give ASP encodings of our formulations in the previous sections.

Encoding the Entities and Events: Rules t1-t2 in the following state that an instance X is an event or an entity if and only if it is the instance of *event*

class or *entity* class respectively. Rules t3-t4 identify E as an event if there is an ordering edge to E (Definition 4.ii). Rule t5 encodes Definition 4.iv. “has(X , ancestorclass, Y)” denotes the transitive closure of “has(M , superclass, N)” and is encoded the standard way (rules t6-t7). The rest of Definition 4 are encoded similarly in rules t8-t21.

```
t1: event(X):-has(X,instance_of,event).
t2: entity(X):-has(X,instance_of,entity).
t3: ordering_edge(next_event; enables; causes; prevents; inhibits).
t4: has(E, instance_of, event) :- has(X, S, E), ordering_edge(S).
t5: has(E, instance_of, event) :- has(X, instance_of, ClassY), has(ClassY,
    ancestorclass, event).
```

Finding Next Events, First Subevents and Last Subevents: Rules e1-e2 find the next events (Definition 5) and rules e3-e6 find the first subevents and the last subevents (Definition 6).

```
e1: predicates(ordering_edge, enables; causes; prevents; inhibits).
e2: has(E1, next_event, E2) :- has(E1, Predicate, E2),
    predicates(ordering_edge, Predicate).
e3: not_fse(Z, E) :- has(Z, subevent, E), has(Z, subevent, E2), E2 != E,
    has(E2, next_event, E).
e4: not_lse(Z, E) :- has(Z, subevent, E), has(Z, subevent, E2), E2 != E,
    has(E, next_event, E2).
e5: has(Z, first_subevent, E) :- has(Z, subevent, E), not not_fse(Z, E).
e6: has(Z, last_subevent, E) :- has(Z, subevent, E), not not_lse(Z, E).
```

Encoding Transport Events and Operational Events: $t_event(E)$ or $o_event(E)$ is used to indicate a transport event or an operational event, respectively.

```
ev1: predicates(t_event, move_through; move_into; move_out_of).
ev2: t_event(E) :- has(E, instance_of, Transport_class), predicates(t_event,
    Transport_class), event(E).
ev3: o_event(E) :- event(E), not t_event(E).
```

Encoding the Inputs and Outputs of Operational Events: We denote the input/output/input location/output location of an event by *input*, *output*, *input_loc* and *output_loc* respectively. Rules i1-i5 get the IOs of operational events. IOs of transport events are encoded similarly (rules i6-i10).

```
i1: input(E,A):-has(E,object,A),o_event(E).
i2: input(E,A):-has(E,base,A),o_event(E).
i3: input(E,A):-has(E,raw_material,A),o_event(E).
i4: output(E,A):-has(E,result,A),o_event(E).
i5: input_loc(E,A):-has(E,site,A),o_event(E).
```

Getting the Missing Inputs and Outputs: Rule i11 gets the input of an event from its first subevent (Definition 7.1). Rule i12 gets the *object* property of a transport event from its first subevent. Other rules to get the input location, output and output location as well as other properties, such as raw-material, result, are encoded in a similar way (rules i13-i24). Rule i25 gets the default output location of an event(Definition 8).

```
i11: has(E, input, A) :- has(SE, input, A), has(E, first_subevent, SE).
i12: has(E, object, A) :- has(SE, object, A), has(E, first_subevent, SE),
    transport_event(E).
i25: has(E, output_location, A) :- not has(E, output_location, A2), has(E,
    input_location, A), entity(A2), event(E), A2 != A.
```

Encoding the Main Class(es) of an Instance: *ClassA* is a main class of instance *A* if *ClassA* is one of *A*'s classes and we do not have *not_main_class(A, ClassA)* (which mean *ClassA* is not the main class of *A*).

```
m1: general_class(thing; event; entity; spatial_entity; tangible_entity;
    chemical_entity).
m2: not_main_class(A, ClassB) :- has(A, instance_of, ClassA), has(A,
    instance_of, ClassB), has(ClassA, ancestorclass, ClassB).
m3: not_main_class(A, ClassB) :- has(A, instance_of, ClassA), has(A,
    instance_of, ClassB), general_class(ClassB), not general_class(ClassA).
m4: main_class(A, ClassA) :- has_class(A, ClassA), not not_main_class(A,
    ClassA).
```

Encoding Instance Matching: We use predicate *match_with(A, B, Confidence)* to represent *match with* relation (Definition 10) from instance *A* to *B*; *Confidence* can be either *low*, *medium* or *high*. Rule ma1 encodes the sub-case 10.1.a of Definition 10. The last rule is for Definition 10.4, matching *A* to *B* transitively through *C*. *lowest_confidence(Conf1, Conf2, Conf)* means *Conf* is the lowest confidence in *Conf1* and *Conf2* (Rules lc1-lc7). Rules for other cases of Definition 10 are skipped (rules ma2-ma5); locational instance matching is encoded in a similar way (rules sma1-sma4).

```
ma1: match_with(A, B, high) :- main_class(A, ClassA), main_class(B, ClassB),
    A==B.
ma6: match_with(A, B, Conf) :- match_with(A, C, Conf1), match_with(C, B, Conf2),
    A!=B, A!=C, B!=C, lowest_confidence(Conf1, Conf2, Conf).
```

Encoding Possible-next-event Relation: In this section, we show how Definition 14 is encoded. We use *has(A, tc_subevent, B)* to represent transitive closure of *sub event* relation between *A* and *B* (encoded by *has(A, subevent, B)*), which is defined in the standard way (rules tsub1-tsub3). We also use *_join(A, B)* to encode that *A* joins to *B* according to Definition 13 (rules j1-j3). The two rules below is corresponding to the sub-case 14.1. Other cases are skipped (rules n2-n5).

```
n1: _notNextEvent(A, B) :- _join(A, SuperB), _join(A, B), has(SuperB,
    tc_subevent, B).
n6: possible_next_event(A, B) :- _join(A, B), not _notNextEvent(A, B).
```

Correctness of the ASP Rules:

Definition 15. The ASP program Π_Z is the answer set program consisting of the facts of the form “*has(X, S, V)*” that are generated from all the nodes and edges of $KDG(Z)$ in the following way:

1. For each node *N*, generate “*has(N, instance_of, event)*” if *N* is event node, “*has(N, instance_of, entity)*” if *N* is entity nodes.
2. For each edge of relation *R* (Table 1) from *E1* to *E2*, generate “*has(E1, R, E2)*”.

Definition 16. The ASP program Π is the answer set program consisting of the following rules: *t1* to *t14* for events and entities, *e1* to *e6* for next events, first subevents and last events, *ev1* to *ev3* for two types of events, *i1* to *i25* for inputs, outputs of events, *m1* to *m4* for main class(es), *lc1* to *lc7* for the lowest confidence, *ma1* to *ma6* for match relation, *sma1* to *sma4* for spatially match relation, *tsub1* to *tsub2* for transitive closure of subevents, *j1* to *j3* for joined events, and *n1* to *n6* for possible next events.

Proposition 1: E is the last subevent of X in $KDG(Z)$ iff

$$\Pi_Z \cup \Pi \models \text{has}(Z, \text{last_subevent}, E)$$

Proposition 2: A is the main class of E in $KDG(Z)$ iff

$$\Pi_Z \cup \Pi \models \text{main_class}(E, A)$$

Proposition 3: Let A and B be two instances in $KDG(Z)$. A matches with B with the confidence level $Conf$ iff $\Pi_Z \cup \Pi \models \text{match_with}(A, B, Conf)$

Proposition 4: Let A and B be two events in $KDG(Z)$. A is a possible next event of B iff $\Pi_Z \cup \Pi \models \text{possible_next_event}(A, B)$

8 Discussion: Answering “How” and “Why” Questions

In Section 4, we showed how to recover missing information using properties of KDG’s structure. Completing this information not only allows us to improve the KB that was used to construct the KDG, but also make it possible to reason about large curated KB using KDG. In Section 5 and 6, we also solved an important step in bringing the KB’s usage out of small examples: we proposed the methods to compare instances and demonstrated their power in finding possible next events.

Those efforts have enabled us to answer deep reasoning questions, such as “How” and “Why” questions. We give examples of a few of them in the following. Details about answering them are explained in another work of ours [2].

1. The answer of “How does X occur?” is simply a structure that basically contains $KDG(X)$ and all the nodes connected to/from X through ordering edges.
2. The answer of “How does X produce Y ?” is similar to “How does X occur?” but X must produce Y .
3. The answer of “How are X and Y related?” is a simplified structure of $KDG(Z)$ that contains: two paths of component edges to X and Y from their lowest common ancestor and all paths of ordering edges linking two nodes in those two paths.
4. Similarly, the answer of “Why X is important to Y ?” is the answer of “How are X and Y related?” plus the path on “important” links which explains why X is important to Y . An “important” link from A to B is defined in KDG to indicate that A is important to B .
5. Other questions that KDG can answer includes “How does X participate in process Y ?”, “How does X do Y ?”, “Why does X produce Y ?” and others.

9 Conclusion

In this paper we have shown how to derive certain missing information from large knowledge bases. Often such knowledge bases are created by multiple people; sometimes even through crowd-sourcing. This often leads to some information being not explicitly stated, even though the knowledge base contains clues to derive that information. In our larger quest to formulate answers to “why” and “how” questions, we focused on the frame based knowledge base AURA, noticed several such omissions, and using those as examples, developed several general formulations regarding missing knowledge about events. We also gave an ASP

implementation of our formulations and used them in answering “why” and “how” questions. We briefly discussed some of those question types and how their answer can be obtained from Knowledge Description Graphs (KDGs). Thus, by being able to obtain missing information and enriching the original KDGs one can obtain more accurate and intuitive answers to the various ‘why’ and ‘how’ questions.

One of our formulations was about entity resolution where we resolve multiple entities that may have different names but may refer to the same entity. Our method is different from other methods in the literature [4, 5]. Since each entity resolution method heavily relies on the properties of the database it is working on, and no other system we know of is about AURA or similar event centered knowledge bases we were unable to directly compare our method with the others.

Our approach to use rules (albeit ASP rules) to derive missing information is analogous to use of rules in data cleaning and in improving data quality [10–12]. However those works do not focus on issues that we discussed in this paper.

References

1. Baral, C., Vo, N.H., Liang, S.: Answering why and how questions with respect to a frame-based knowledge base: a preliminary report. Technical Communications of the 28th International Conference on Logic Programming (ICLP’12) **17** (2012) 26–36
2. Baral, C., Vo, N.: Formulating question answering with respect to event-object description graphs. Unpublished paper submitted to a conference (2013)
3. Chaudhri, V.K., Clark, P.E., Mishra, S., Pacheco, J., Spaulding, A., Tien, J.: AURA: capturing knowledge and answering questions on science textbooks. Technical report, SRI International (2009)
4. Getoor, L., Diehl, C.P.: Link mining: a survey. ACM SIGKDD Explorations Newsletter **7**(2) (2005) 3–12
5. Brizan, D.G., Tansel, A.U.: A survey of entity resolution and record linkage methodologies. Communications of the IIMA **6**(3) (2006) 41–50
6. Baral, C., Liang, S.: From knowledge represented in frame-based languages to declarative representation and reasoning via ASP. 13th International Conference on Principles of Knowledge Representation and Reasoning (2012)
7. Gelfond, M., Lifschitz, V.: The stable model semantics for logic programming. In Kowalski, R., Bowen, K., eds.: Logic Programming: Proc. of the Fifth Int’l Conf. and Symp., MIT Press (1988) 1070–1080
8. Baral, C.: Knowledge representation, reasoning and declarative problem solving. Cambridge University Press (2003)
9. Clark, P., Porter, B., Works, B.: KM: The knowledge machine 2.0: Users manual. Citeseer (2004)
10. Herzog, T.N., Scheuren, F.J., Winkler, W.E.: Data quality and record linkage techniques. Springer (2007)
11. Rahm, E., Do, H.H.: Data cleaning: Problems and current approaches. IEEE Data Engineering Bulletin **23**(4) (2000) 3–13
12. Fan, W., Geerts, F., Jia, X.: Conditional dependencies: A principled approach to improving data quality. In: Dataspace: The Final Frontier. Springer (2009) 8–20